

## Importing all necessary libraries

```
In [ ]: from IPython.display import Audio, display
import IPython.display as ipd

import librosa
from librosa import display

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

import tensorflow as tf
from tensorflow.keras.callbacks import EarlyStopping

from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.preprocessing import StandardScaler

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv1D, MaxPooling1D, Flatten, Dense
from tensorflow.keras.utils import to_categorical

from tensorflow.keras.models import load_model
```

## Function for loading dataset

```
In [ ]: def load_dataset(path):
        """Loads dataset from path provided and return pandas object"""
        ds_path= path
        ds= pd.read_csv(ds_path)
        return ds
```

## Function for loading audios using librosa

```
In [ ]: def audio_signals(path, sample_rate=16000, duration=3):
        """Loading audio using librosa library"""
        audio, sr= librosa.load(path, sr=sample_rate, duration=duration, mono=True)
        if sr!=sample_rate:
            audio = librosa.resample(audio, sr, sample_rate)
        return audio

def load_audio_signals(ds, num_of_samples):
    """Returns numpy array of all audio signals"""
    all_audio_signals=[]

    for path in ds['file_path']:
        audio= audio_signals("./Dataset-2/"+path)
        all_audio_signals.append(audio)
        if(len(all_audio_signals)==num_of_samples):
            break

    all_audio_signals= np.array(all_audio_signals, dtype=object)
    return all_audio_signals
```

## Data augmentation (Data Pre-processing)

```

In [ ]: import numpy as np
import librosa
import random

def add_noise(audio, noise_factor=0.005):
    """Add random noise to the audio signal"""
    noise = np.random.randn(len(audio))
    augmented_audio = audio + noise_factor * noise
    return augmented_audio

def time_stretch(audio, rate=1.2):
    """Stretch or compress the audio in time without changing pitch"""
    augmented_audio = librosa.effects.time_stretch(y=audio, rate=rate)
    return augmented_audio

def pitch_shift(audio, sr, n_steps=2):
    """Shift the pitch of the audio signal"""
    augmented_audio = librosa.effects.pitch_shift(audio, sr=sr, n_steps=n_steps)
    return augmented_audio

def time_shift(audio, sr, shift_factor=0.2):
    """Shift the audio signal in time"""
    shift_samples = int(shift_factor * len(audio))
    augmented_audio = np.roll(audio, shift_samples)
    return augmented_audio

def apply_augmentation(audio, sr=16000):
    """Apply random augmentation to the audio signal"""
    # augmentation_functions = [add_noise, time_stretch, pitch_shift, time_shift]
    # augmentation_functions = [add_noise, time_stretch, pitch_shift, lambda x: time_shift(x, sr, 0.1)]
    augmentation_functions = [add_noise, time_stretch, lambda x, sr=sr: pitch_shift(x, sr, 0.1)]
    augmentation_function = random.choice(augmentation_functions)
    augmented_audio = augmentation_function(audio)
    return augmented_audio

def augment_audios(audios):
    """Apply augmentation to a list of audio signals"""
    augmented_audios = []
    for audio in audios:
        augmented_audio = apply_augmentation(audio)
        augmented_audios.append(augmented_audio)
    return np.array(augmented_audios, dtype=object)

```

## Label encoding for speakers

```

In [ ]: def label_encoding(ds, num_of_samples):
    """Label encoding of speakers for classification"""
    labels=[]

    for label in ds['speaker']:
        labels.append(label)

    # Initialize the LabelEncoder
    label_encoder = LabelEncoder()

    # Fit and transform the labels
    y_encoded = label_encoder.fit_transform(labels)

    # Print the mapping between original labels and encoded values
    label_mapping = dict(zip(label_encoder.classes_, label_encoder.transform(labels)))
    print("Label Mapping:", label_mapping)

```

```
y_encoded=y_encoded[0:num_of_samples]
return y_encoded
```

## Fourier Magnitude Spectra- Feature Extraction

```
In [ ]: def feature_extraction(all_audio_signals, num_of_samples):
        """Returns the fourier magnitude spectrum as feature vector"""
        fourier_magnitude=[]
        fixed_num_freq_samples = 24001

        for audio in all_audio_signals[0:num_of_samples]:
            fourier = np.fft.rfft(audio)
            sampling_rate = 16000.0

            num_freq_samples = len(fourier)
            resampled_fourier = np.interp(np.linspace(0, num_freq_samples - 1, fixed_num_freq_samples),
            np.arange(0, num_freq_samples), fourier)
            N = len(audio)
            normalize = N/2
            norm_amplitude = np.abs(resampled_fourier)/normalize
            fourier_magnitude.append(norm_amplitude)

        fourier_magnitude= np.array(fourier_magnitude, dtype=float)

        return fourier_magnitude
```

## For Training

```
In [ ]: ds= load_dataset("./Dataset-2/train.csv")
```

```
In [ ]: num_of_samples= 3000
audios= load_audio_signals(ds, num_of_samples)
encoded_y= label_encoding(ds,num_of_samples)
print(audios)
print(audios.shape)

Label Mapping: {'aew': 0, 'ahw': 1, 'aup': 2, 'awb': 3, 'axb': 4, 'bd1': 5, 'clb': 6, 'eey': 7, 'fem': 8, 'gka': 9, 'jmk': 10, 'ksp': 11, 'ljm': 12, 'lnh': 13, 'rms': 14, 'rxr': 15, 'slp': 16, 'slt': 17}
[array([0.00521851, 0.0062561 , 0.00543213, ..., 0.          , 0.          ,
        0.          ], dtype=float32)
 array([0.00094604, 0.00115967, 0.00067139, ..., 0.01141357, 0.01730347,
        0.00491333], dtype=float32)
 array([-0.00115967, -0.00143433, -0.00137329, ..., -0.00183105,
        -0.00195312, -0.00183105], dtype=float32)
 array([ 0.00012207,  0.00012207,  0.00039673, ...,  0.06216431,
         0.0015564 , -0.02697754], dtype=float32)
 array([-0.0005188 , -0.0005188 , -0.00048828, ...,  0.          ,
         0.          ,  0.          ], dtype=float32)
 array([-0.00430298, -0.00424194, -0.00469971, ..., -0.0007019 ,
        -0.00048828, -0.0007019 ], dtype=float32)
(3000,)
```

```
In [ ]: augmented_audios = augment_audios(audios)
print(augmented_audios)
print(augmented_audios.shape)
```

```
[array([-0.05493164, -0.05148315, -0.04650879, ..., -0.05032349,
        -0.05447388, -0.05737305], dtype=float32)
 array([ 0.00094406,  0.00106522,  0.00078033, ..., -0.00802859,
        -0.00226684,  0.00577904], dtype=float32)
 array([ 0.00704814, -0.00129165, -0.00388875, ..., -0.00016747,
        -0.00443582, -0.00039775])
 ...,
 array([7.3119605e-05, 1.9680563e-04, 3.4200089e-04, ..., 2.4078891e-03,
        5.7600420e-03, 3.1722693e-03], dtype=float32)
 array([-0.0004702 , -0.00052782, -0.00049365, ..., -0.00030277,
        -0.00027932, -0.00032359], dtype=float32)
 array([-4.2791897e-03, -4.2221909e-03, -4.6514133e-03, ...,
        1.2576625e-04,  2.8581971e-05, -3.3309010e-05], dtype=float32)]
(3000,)
```

```
In [ ]: features= feature_extraction(augmented_audios,num_of_samples)
```

## For Testing

```
In [ ]: test_ds= load_dataset("./Dataset-2/test_full.csv")
test_audios= load_audio_signals(test_ds,num_of_samples)
encoded_y_test= label_encoding(test_ds,num_of_samples)
features_y= feature_extraction(test_audios,num_of_samples)
```


Label Mapping: {'aew': 0, 'ahw': 1, 'aup': 2, 'awb': 3, 'axb': 4, 'bd1': 5, 'clb': 6, 'eey': 7, 'fem': 8, 'gka': 9, 'jmk': 10, 'ksp': 11, 'ljm': 12, 'lnh': 13, 'rm s': 14, 'rxr': 15, 'slp': 16, 'slt': 17}

## Plots of augmented audios

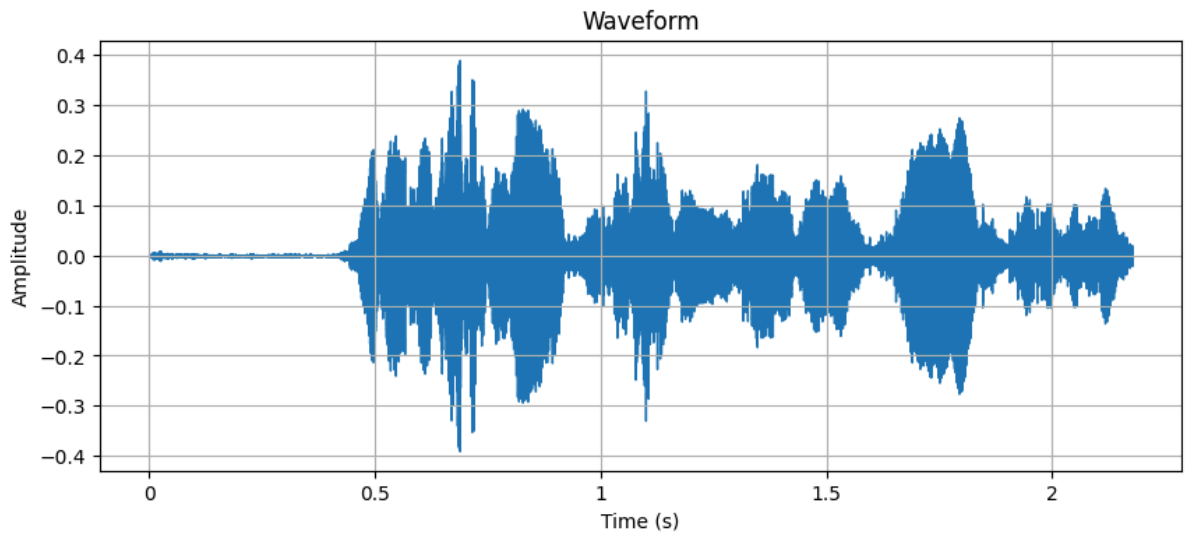
```
In [ ]: # plotting first 2 signals

for i in range(2):
    plt.figure()
    #plot the raw audio signal using librosa
    plt.figure(figsize=(10,4))
    librosa.display.waveshow(augmented_audios[i])
    plt.xlabel("Time (s)")
    plt.ylabel("Amplitude")
    plt.title("Waveform")
    plt.grid()
    plt.show()
```

<Figure size 640x480 with 0 Axes>



<Figure size 640x480 with 0 Axes>




```
In [ ]: for i in range(2):
    print(f"For speaker {i}")
    fourier = np.fft.rfft(augmented_audios[i])
    # Get the frequency components of the spectrum
    sampling_rate = 16000.0 # It's used as a sample spacing
    frequency_axis = np.linspace(0, sampling_rate, len(np.abs(fourier)))
    plt.figure()
    # Plot the result (the spectrum |Xk|)
    plt.plot(frequency_axis, np.abs(fourier))
    plt.show()


    # Calculate N/2 to normalize the FFT output
    N = len(audios[i])
    normalize = N/2

    plt.figure()
    # Plot the normalized FFT (|Xk|)/(N/2)
    plt.plot(frequency_axis, np.abs(fourier)/normalize)
    plt.ylabel('Amplitude')
    plt.xlabel('Samples')
    plt.title('Normalized FFT Spectrum')
    plt.show()
```


For speaker 0




Normalized FFT Spectrum



For speaker 1



Normalized FFT Spectrum



## Machine Learning: Logistic Regression

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(features, encoded_y, test_size
```

```
In [ ]: sc = StandardScaler()  
X_train = sc.fit_transform(X_train)  
X_test = sc.transform(X_test)
```

```
In [ ]: classifier = LogisticRegression(random_state = 0)
classifier.fit(X_train, y_train)
```

```
Out[ ]: LogisticRegression
LogisticRegression(random_state=0)
```

## TESTING

```
In [ ]: print(classifier.predict(sc.transform(features_y)))
y_pred=classifier.predict(sc.transform(features_y))

[ 7  7  3 ... 11  7 15]
```

```
In [ ]: y_test=encoded_y_test
print(y_test)

[ 0  7  3 ... 11  7 15]
```

```
In [ ]: cm = confusion_matrix(y_test, y_pred)
print(cm)
print("Accuracy", accuracy_score(y_pred, y_test))

[[172  4  2  1  7  3  1  6  5  0  1  2  0  4  4  4  2  1]
 [  0 106  0  0  0  0  0  0  1  0  0  4  0  0  0  2  0  0]
 [  0  0 112  0  4  1  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0 213  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0 111  0  0  1  0  0  0  0  0  1  0  0  1  1]
 [  1  1  3  7  1 195  0  2  1  0  2  1  0  3  0  0  0  0]
 [  0  0  0  0  0  1 211  0  0  0  0  0  0  0  0  0  0  8]
 [  0  1  0  0  1  3  1  90  0  0  0  0  3  2  0  0  12  1]
 [  0  0  0  0  1  0  0  0 113  0  0  0  0  0  0  0  0  0]
 [ 16  0  0  0  0  0  0  0  0  91  2  5  0  0  0  1  0  0]
 [  0  0  0  0  0  0  0  0  0  0 218  0  0  0  0  0  0  0]
 [  1  0  0  0  0  0  0  0  0  0  0 217  0  0  0  0  1  0]
 [  0  0  0  0  3  0  2  3  0  0  0  0  78  8  0  0  18  3]
 [  1  0  0  0  0  8  0  0  0  0  0  0  3 201  1  0  2  2]
 [  4  0  1  0  0  0  0  1  1  0  0  0  0  0 205  0  0  0]
 [  1  6  3  0  1  0  0  0  1  1  0  6  0  4  0 104  1  1]
 [  0  0  0  0  3  0  2  3  0  0  0  0  6  0  0  0 100  0]
 [  0  2  0  1  0  0  2  0  1  0  3  0  1  0  0  0  0 208]]

Accuracy 0.915
```

```
In [ ]: from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



	precision	recall	f1-score	support
0	0.88	0.79	0.83	219
1	0.88	0.94	0.91	113
2	0.93	0.96	0.94	117
3	0.96	1.00	0.98	213
4	0.84	0.97	0.90	115
5	0.92	0.90	0.91	217
6	0.96	0.96	0.96	220
7	0.85	0.79	0.82	114
8	0.92	0.99	0.95	114
9	0.99	0.79	0.88	115
10	0.96	1.00	0.98	218
11	0.92	0.99	0.96	219
12	0.86	0.68	0.76	115
13	0.90	0.92	0.91	218
14	0.98	0.97	0.97	212
15	0.94	0.81	0.87	129
16	0.73	0.88	0.80	114
17	0.92	0.95	0.94	218
accuracy			0.92	3000
macro avg	0.91	0.90	0.90	3000
weighted avg	0.92	0.92	0.91	3000

```
In [ ]: from sklearn.metrics import precision_recall_fscore_support
res = []
for l in range(18):
    prec, recall, _, _ = precision_recall_fscore_support(np.array(y_test)==l,
                                                         np.array(y_pred)==l,
                                                         pos_label=True, average=None)
    res.append([l, recall[0], recall[1]])

pd.DataFrame(res, columns = ['class', 'specificity', 'sensitivity'])
```

Out[ ]:

	class	specificity	sensitivity
0	0	0.991370	0.785388
1	1	0.995151	0.938053
2	2	0.996878	0.957265
3	3	0.996771	1.000000
4	4	0.992721	0.965217
5	5	0.994251	0.898618
6	6	0.997122	0.959091
7	7	0.994456	0.789474
8	8	0.996535	0.991228
9	9	0.999653	0.791304
10	10	0.997124	1.000000
11	11	0.993528	0.990868
12	12	0.995494	0.678261
13	13	0.992092	0.922018
14	14	0.998207	0.966981
15	15	0.997562	0.806202
16	16	0.987179	0.877193
17	17	0.993889	0.954128

## 1D CNN without Windowing

```
In [ ]: X_train, X_temp, y_train, y_temp = train_test_split(features, encoded_y, test_size=
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, rand
```

```
X_train = X_train.astype('float32')
y_train = y_train.astype('int32')
X_val = X_val.astype('float32')
y_val = y_val.astype('int32')
X_test = X_test.astype('float32')
y_test = y_test.astype('int32')
```

```
# Print the shapes of training and validation data
```

```
print("Training Data Shape:", X_train.shape)
```

```
print("Validation Data Shape:", X_val.shape)
```

```
print("Test Data Shape:", X_test.shape)
```

```
print("Training y Data Shape:", y_train.shape)
```

```
print("Validation y Data Shape:", y_val.shape)
```

```
print("Test y Data Shape:", y_test.shape)
```

```
Training Data Shape: (2100, 24001)
```

```
Validation Data Shape: (450, 24001)
```

```
Test Data Shape: (450, 24001)
```

```
Training y Data Shape: (2100,)
```

```
Validation y Data Shape: (450,)
```

```
Test y Data Shape: (450,)
```

```

In [ ]: # Reshape the input data for compatibility with Conv1D
X_train_reshaped = X_train.reshape(X_train.shape[0], X_train.shape[1], 1)
X_val_reshaped = X_val.reshape(X_val.shape[0], X_val.shape[1], 1)

# Convert labels to one-hot encoded format
num_classes = len(np.unique(y_train))
y_train_encoded = to_categorical(y_train, num_classes=num_classes)
y_val_encoded = to_categorical(y_val, num_classes=num_classes)

# Defining the CNN model
model = Sequential([
    Conv1D(64, kernel_size=3, activation='relu', input_shape=(X_train.shape[1], 1)),
    MaxPooling1D(pool_size=2),
    Conv1D(128, kernel_size=3, activation='relu'),
    MaxPooling1D(pool_size=2),
    Flatten(),
    Dense(64, activation='relu'),
    Dense(num_classes, activation='softmax')
])

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model.summary()

history = model.fit(X_train_reshaped, y_train_encoded, validation_data=(X_val_reshaped, y_val_encoded))

model.save("1d_cnn_model_without_window_new.h5")

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv1d_2 (Conv1D)	(None, 23999, 64)	256
max_pooling1d_2 (MaxPooling1D)	(None, 11999, 64)	0
conv1d_3 (Conv1D)	(None, 11997, 128)	24704
max_pooling1d_3 (MaxPooling1D)	(None, 5998, 128)	0
flatten_1 (Flatten)	(None, 767744)	0
dense_2 (Dense)	(None, 64)	49135680
dense_3 (Dense)	(None, 18)	1170
Total params: 49,161,810		
Trainable params: 49,161,810		
Non-trainable params: 0		

Layer (type)	Output Shape	Param #
conv1d_2 (Conv1D)	(None, 23999, 64)	256
max_pooling1d_2 (MaxPooling1D)	(None, 11999, 64)	0
conv1d_3 (Conv1D)	(None, 11997, 128)	24704
max_pooling1d_3 (MaxPooling1D)	(None, 5998, 128)	0
flatten_1 (Flatten)	(None, 767744)	0
dense_2 (Dense)	(None, 64)	49135680
dense_3 (Dense)	(None, 18)	1170
Total params: 49,161,810		
Trainable params: 49,161,810		
Non-trainable params: 0		

Epoch 1/20  
66/66 [=====] - 151s 2s/step - loss: 2.8230 - accuracy: 0.1181 - val\_loss: 2.5752 - val\_accuracy: 0.2200  
Epoch 2/20  
66/66 [=====] - 135s 2s/step - loss: 1.9129 - accuracy: 0.4086 - val\_loss: 1.3391 - val\_accuracy: 0.5333  
Epoch 3/20  
66/66 [=====] - 134s 2s/step - loss: 0.9789 - accuracy: 0.6738 - val\_loss: 0.8625 - val\_accuracy: 0.6644  
Epoch 4/20  
66/66 [=====] - 142s 2s/step - loss: 0.6184 - accuracy: 0.7957 - val\_loss: 0.6828 - val\_accuracy: 0.7800  
Epoch 5/20  
66/66 [=====] - 145s 2s/step - loss: 0.4659 - accuracy: 0.8471 - val\_loss: 0.4657 - val\_accuracy: 0.8467

```

Epoch 6/20
66/66 [=====] - 156s 2s/step - loss: 0.3950 - accuracy:
0.8700 - val_loss: 0.4366 - val_accuracy: 0.8422
Epoch 7/20
66/66 [=====] - 142s 2s/step - loss: 0.2873 - accuracy:
0.9043 - val_loss: 0.3922 - val_accuracy: 0.8800
Epoch 8/20
66/66 [=====] - 146s 2s/step - loss: 0.1939 - accuracy:
0.9424 - val_loss: 0.3715 - val_accuracy: 0.8889
Epoch 9/20
66/66 [=====] - 145s 2s/step - loss: 0.1778 - accuracy:
0.9424 - val_loss: 0.3822 - val_accuracy: 0.8756
Epoch 10/20
66/66 [=====] - 143s 2s/step - loss: 0.1463 - accuracy:
0.9476 - val_loss: 0.3227 - val_accuracy: 0.9044
Epoch 11/20
66/66 [=====] - 139s 2s/step - loss: 0.1041 - accuracy:
0.9676 - val_loss: 0.4096 - val_accuracy: 0.8622
Epoch 12/20
66/66 [=====] - 135s 2s/step - loss: 0.0824 - accuracy:
0.9752 - val_loss: 0.2775 - val_accuracy: 0.9178
Epoch 13/20
66/66 [=====] - 132s 2s/step - loss: 0.0901 - accuracy:
0.9719 - val_loss: 0.3343 - val_accuracy: 0.9133
Epoch 14/20
66/66 [=====] - 132s 2s/step - loss: 0.0669 - accuracy:
0.9805 - val_loss: 0.3166 - val_accuracy: 0.9222
Epoch 15/20
66/66 [=====] - 157s 24s/step - loss: 0.0368 - accuracy:
0.9919 - val_loss: 0.2981 - val_accuracy: 0.9222
Epoch 16/20
66/66 [=====] - 154s 2s/step - loss: 0.0160 - accuracy:
0.9995 - val_loss: 0.3055 - val_accuracy: 0.9244
Epoch 17/20
66/66 [=====] - 146s 2s/step - loss: 0.0117 - accuracy:
0.9990 - val_loss: 0.3093 - val_accuracy: 0.9178
Epoch 18/20
66/66 [=====] - 141s 2s/step - loss: 0.0082 - accuracy:
1.0000 - val_loss: 0.2933 - val_accuracy: 0.9244
Epoch 19/20
66/66 [=====] - 142s 2s/step - loss: 0.0062 - accuracy:
1.0000 - val_loss: 0.3050 - val_accuracy: 0.9222
Epoch 20/20
66/66 [=====] - 130s 2s/step - loss: 0.0058 - accuracy:
1.0000 - val_loss: 0.3447 - val_accuracy: 0.9222



```

```

In [ ]: # Plot training & validation accuracy values
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Plot training & validation loss values
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

```



## TESTING

```
In [ ]: loaded_model = load_model("1d_cnn_model_without_window_new.h5")  
  
X_test_resaped = features_y.reshape(features_y.shape[0], features_y.shape[1], 1)
```

```
y_pred = loaded_model.predict(X_test_resaped)
```

```
94/94 [=====] - 46s 491ms/step
```

```
In [ ]: y_pred_classes = np.argmax(y_pred, axis=1)
print(y_pred_classes)
```

```
[ 0  7  3 ... 11  7 15]
```

```
In [ ]: from sklearn.metrics import confusion_matrix, accuracy_score
cm = confusion_matrix(encoded_y_test, y_pred_classes)
print(cm)
accuracy_score(encoded_y_test, y_pred_classes)
```

```
[[193  0  3  0  0  0  0  0  6  2  1  6  0  0  8  0  0  0]
 [  0 105  0  1  0  0  0  0  0  0  0  7  0  0  0  0  0  0]
 [  1  0 113  1  0  1  0  0  0  0  1  0  0  0  0  0  0  0]
 [  0  0  0 213  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  2  0 91  0  0  3  0  0  0  0  8  2  0  0  6  3]
 [  0  1 14  6  0 182  0  0 12  0  0  0  0  2  0  0  0  0]
 [  0  0  0  0  0  0 212  0  0  0  0  0  0  0  0  0  0  8]
 [  0  0  0  0  0  0  3 96  0  0  0  0  2  0  0  0  8  5]
 [  0  0  0  0  0  0  0  0 106  0  0  0  0  0  8  0  0  0]
 [  0  0  0  0  0  0  0  0  0 106  8  0  0  0  1  0  0  0]
 [  0  0  0  0  0  0  0  0  0  0 217  0  0  0  1  0  0  0]
 [  3 10  0  0  0  3  0  0  2  5  1 195  0  0  0  0  0  0]
 [  0  0  0  0  2  0  0 10  0  0  0  0 66 14  0  1 13  9]
 [  1  0  3  0  0  0  0  2  0  0  0  0  0 206  0  3  2  1]
 [  1  0  0  0  0  0  0  0  0  0  0  0  0  0 211  0  0  0]
 [  0  6  0  0  0  0  0  0  0  1  0  4  0  2  0 116  0  0]
 [  0  0  0  0  2  0  3 25  0  0  0  0  9  0  0  0 72  3]
 [  0  0  2  2  0  0  2  1  0  0  0  6  1  0  0  1  0 203]]
```

```
Out[ ]: 0.901
```

```
In [ ]: from sklearn.metrics import classification_report
print(classification_report(encoded_y_test,y_pred_classes))
```

	precision	recall	f1-score	support
0	0.97	0.88	0.92	219
1	0.86	0.93	0.89	113
2	0.82	0.97	0.89	117
3	0.96	1.00	0.98	213
4	0.96	0.79	0.87	115
5	0.98	0.84	0.90	217
6	0.96	0.96	0.96	220
7	0.70	0.84	0.76	114
8	0.84	0.93	0.88	114
9	0.93	0.92	0.93	115
10	0.95	1.00	0.97	218
11	0.89	0.89	0.89	219
12	0.77	0.57	0.66	115
13	0.91	0.94	0.93	218
14	0.92	1.00	0.96	212
15	0.96	0.90	0.93	129
16	0.71	0.63	0.67	114
17	0.88	0.93	0.90	218
accuracy			0.90	3000
macro avg	0.89	0.88	0.88	3000
weighted avg	0.90	0.90	0.90	3000

```
In [ ]: from sklearn.metrics import precision_recall_fscore_support
res = []
for l in range(18):
    prec, recall, _, _ = precision_recall_fscore_support(np.array(encoded_y_test)==l,
                                                         np.array(y_pred_classes)==l,
                                                         pos_label=True, average=None)
    res.append([l, recall[0], recall[1]])

pd.DataFrame(res, columns = ['class', 'specificity', 'sensitivity'])
```

```
Out[ ]:
```

	class	specificity	sensitivity
0	0	0.997843	0.881279
1	1	0.994112	0.929204
2	2	0.991675	0.965812
3	3	0.996412	1.000000
4	4	0.998614	0.791304
5	5	0.998563	0.838710
6	6	0.997122	0.963636
7	7	0.985793	0.842105
8	8	0.993070	0.929825
9	9	0.997227	0.921739
10	10	0.996046	0.995413
11	11	0.991730	0.890411
12	12	0.993068	0.573913
13	13	0.992811	0.944954
14	14	0.993544	0.995283
15	15	0.998258	0.899225
16	16	0.989951	0.631579
17	17	0.989576	0.931193

## 1D CNN with hamming window

```
In [ ]: def apply_window(audio, window_type="hamming"):
    if window_type == "hamming":
        window = np.hamming(len(audio))
    else:
        raise ValueError("Invalid window type. Supported type is 'hamming'.")

    # Apply the window to the audio signal
    windowed_audio = audio * window
    return windowed_audio
```


```
In [ ]: augmented_audios=[apply_window(audio, window_type="hamming") for audio in augmented_audios]
features= feature_extraction(augmented_audios,num_of_samples)
```

## Plots with windowed audios




```
In [ ]: for i in range(2):
plt.figure()
#plot the raw audio signal using librosa
plt.figure(figsize=(10,4))
librosa.display.waveshow(augmented_audios[i])
plt.xlabel("Time (s)")
plt.ylabel("Amplitude")
plt.title("Waveform")
plt.grid()
plt.show()
```

<Figure size 640x480 with 0 Axes>



<Figure size 640x480 with 0 Axes>




```
In [ ]: for i in range(2):
print(f"For speaker {i}")
fourier = np.fft.rfft(augmented_audios[i])
# Get the frequency components of the spectrum
sampling_rate = 16000.0 # It's used as a sample spacing
frequency_axis = np.linspace(0, sampling_rate, len(np.abs(fourier)))
plt.figure()
# Plot the result (the spectrum |Xk|)
plt.plot(frequency_axis,np.abs(fourier))
plt.show()


# Calculate N/2 to normalize the FFT output
N = len(audios[i])
normalize = N/2
```

```
plt.figure()
# Plot the normalized FFT ( $|X_k|/(N/2)$ )
plt.plot(frequency_axis, np.abs(fourier)/normalize)
plt.ylabel('Amplitude')
plt.xlabel('Samples')
plt.title('Normalized FFT Spectrum')
plt.show()
```


For speaker 0




Normalized FFT Spectrum



For speaker 1



Normalized FFT Spectrum



```
In [ ]: # Split the data into training, validation, and test sets
X_train, X_temp, y_train, y_temp = train_test_split(features, encoded_y, test_size=
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, rand

X_train = X_train.astype('float32')
y_train = y_train.astype('int32')
X_val = X_val.astype('float32')
y_val = y_val.astype('int32')
```

```

X_test = X_test.astype('float32')
y_test = y_test.astype('int32')

# Print the shapes of training and validation data
print("Training Data Shape:", X_train.shape)
print("Validation Data Shape:", X_val.shape)
print("Test Data Shape:", X_test.shape)

print("Training y Data Shape:", y_train.shape)
print("Validation y Data Shape:", y_val.shape)
print("Test y Data Shape:", y_test.shape)

```

```

Training Data Shape: (2100, 24001)
Validation Data Shape: (450, 24001)
Test Data Shape: (450, 24001)
Training y Data Shape: (2100,)
Validation y Data Shape: (450,)
Test y Data Shape: (450,)

```

```

In [ ]: # Reshape the input data for compatibility with Conv1D
X_train_resaped = X_train.reshape(X_train.shape[0], X_train.shape[1], 1)
X_val_resaped = X_val.reshape(X_val.shape[0], X_val.shape[1], 1)

# Convert Labels to one-hot encoded format
num_classes = len(np.unique(y_train))
y_train_encoded = to_categorical(y_train, num_classes=num_classes)
y_val_encoded = to_categorical(y_val, num_classes=num_classes)

# Defining the CNN model
model = Sequential([
    Conv1D(64, kernel_size=3, activation='relu', input_shape=(X_train.shape[1], 1)),
    MaxPooling1D(pool_size=2),
    Conv1D(128, kernel_size=3, activation='relu'),
    MaxPooling1D(pool_size=2),
    Flatten(),
    Dense(64, activation='relu'),
    Dense(num_classes, activation='softmax')
])

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model.summary()

history = model.fit(X_train_resaped, y_train_encoded, validation_data=(X_val_resaped, y_val_encoded))

model.save("1d_cnn_model_hamming_window_new.h5")

```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
conv1d_10 (Conv1D)	(None, 23999, 64)	256
conv1d_10 (Conv1D)	(None, 23999, 64)	256
max_pooling1d_10 (MaxPooling1D)	(None, 11999, 64)	0
conv1d_11 (Conv1D)	(None, 11997, 128)	24704
max_pooling1d_11 (MaxPooling1D)	(None, 5998, 128)	0
flatten_5 (Flatten)	(None, 767744)	0
dense_10 (Dense)	(None, 64)	49135680
dense_11 (Dense)	(None, 18)	1170

Total params: 49,161,810  
Trainable params: 49,161,810  
Non-trainable params: 0

Epoch 1/25  
66/66 [=====] - 190s 3s/step - loss: 2.8711 - accuracy: 0.0700 - val\_loss: 2.8470 - val\_accuracy: 0.0689  
Epoch 2/25  
66/66 [=====] - 168s 3s/step - loss: 2.8531 - accuracy: 0.0705 - val\_loss: 2.8466 - val\_accuracy: 0.0711  
Epoch 3/25  
66/66 [=====] - 173s 3s/step - loss: 2.8454 - accuracy: 0.0814 - val\_loss: 2.8320 - val\_accuracy: 0.1444  
Epoch 4/25  
66/66 [=====] - 179s 3s/step - loss: 2.6823 - accuracy: 0.1705 - val\_loss: 2.3973 - val\_accuracy: 0.2800  
Epoch 5/25  
66/66 [=====] - 174s 3s/step - loss: 1.7520 - accuracy: 0.4686 - val\_loss: 1.3742 - val\_accuracy: 0.5711  
Epoch 6/25  
66/66 [=====] - 167s 3s/step - loss: 1.2066 - accuracy: 0.6238 - val\_loss: 1.2749 - val\_accuracy: 0.5756  
Epoch 7/25  
66/66 [=====] - 156s 2s/step - loss: 0.8624 - accuracy: 0.7233 - val\_loss: 0.9356 - val\_accuracy: 0.6822  
Epoch 8/25  
66/66 [=====] - 141s 2s/step - loss: 0.7111 - accuracy: 0.7762 - val\_loss: 0.7188 - val\_accuracy: 0.7644  
Epoch 9/25  
66/66 [=====] - 141s 2s/step - loss: 0.5716 - accuracy: 0.8162 - val\_loss: 0.7168 - val\_accuracy: 0.7556  
Epoch 10/25  
66/66 [=====] - 139s 2s/step - loss: 0.4749 - accuracy: 0.8524 - val\_loss: 0.7868 - val\_accuracy: 0.7333  
Epoch 11/25  
66/66 [=====] - 140s 2s/step - loss: 0.4106 - accuracy: 0.8629 - val\_loss: 0.5235 - val\_accuracy: 0.8422  
Epoch 12/25  
66/66 [=====] - 141s 2s/step - loss: 0.3372 - accuracy:

```



0.8910 - val_loss: 0.5334 - val_accuracy: 0.8089
Epoch 13/25
66/66 [=====] - 141s 2s/step - loss: 0.2641 - accuracy:
0.9248 - val_loss: 0.4321 - val_accuracy: 0.8533
Epoch 14/25
66/66 [=====] - 141s 2s/step - loss: 0.2650 - accuracy:
0.9162 - val_loss: 0.5273 - val_accuracy: 0.8289
Epoch 15/25
66/66 [=====] - 139s 2s/step - loss: 0.1946 - accuracy:
0.9419 - val_loss: 0.4189 - val_accuracy: 0.8689
Epoch 16/25
66/66 [=====] - 140s 2s/step - loss: 0.1492 - accuracy:
0.9557 - val_loss: 0.4318 - val_accuracy: 0.8600
Epoch 17/25
66/66 [=====] - 139s 2s/step - loss: 0.1355 - accuracy:
0.9576 - val_loss: 0.4981 - val_accuracy: 0.8422
Epoch 18/25
66/66 [=====] - 139s 2s/step - loss: 0.1208 - accuracy:
0.9681 - val_loss: 0.5594 - val_accuracy: 0.8378
Epoch 19/25
66/66 [=====] - 140s 2s/step - loss: 0.0958 - accuracy:
0.9743 - val_loss: 0.4734 - val_accuracy: 0.8489
Epoch 20/25
66/66 [=====] - 140s 2s/step - loss: 0.0914 - accuracy:
0.9748 - val_loss: 0.4158 - val_accuracy: 0.8689
Epoch 21/25
66/66 [=====] - 140s 2s/step - loss: 0.0597 - accuracy:
0.9881 - val_loss: 0.5404 - val_accuracy: 0.8289
Epoch 22/25
66/66 [=====] - 140s 2s/step - loss: 0.0468 - accuracy:
0.9929 - val_loss: 0.4779 - val_accuracy: 0.8644
Epoch 23/25
66/66 [=====] - 141s 2s/step - loss: 0.1601 - accuracy:
0.9514 - val_loss: 0.5095 - val_accuracy: 0.8444
Epoch 24/25
66/66 [=====] - 140s 2s/step - loss: 0.0795 - accuracy:
0.9805 - val_loss: 0.5467 - val_accuracy: 0.8422
Epoch 25/25
66/66 [=====] - 140s 2s/step - loss: 0.0667 - accuracy:
0.9819 - val_loss: 0.5163 - val_accuracy: 0.8444

```

```
In [ ]: model.save("1d_cnn_model_hamming_window_new.h5")
```

```
In [ ]: # Plot training & validation accuracy values
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Plot training & validation loss values
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```



## TESTING

```
In [ ]: test_audios=[apply_window(audio, window_type="hamming") for audio in test_audios]
         features_y= feature_extraction(test_audios,num_of_samples)
```

```
In [ ]: loaded_model = load_model("1d_cnn_model_hamming_window_new.h5")

X_test_resaped = features_y.reshape(features_y.shape[0], features_y.shape[1], 1)

y_pred = loaded_model.predict(X_test_resaped)
```

94/94 [=====] - 40s 422ms/step

```
In [ ]: y_pred_classes = np.argmax(y_pred, axis=1)
print(y_pred_classes)
```

[ 0 7 3 ... 11 16 15]

```
In [ ]: from sklearn.metrics import confusion_matrix, accuracy_score
cm = confusion_matrix(encoded_y_test, y_pred_classes)
print(cm)
print("Accuracy:", accuracy_score(encoded_y_test, y_pred_classes))
```

```
[[194  0  2  1  0  0  0  0  8  4  1  6  0  0  3  0  0  0]
 [  0 84  0  3  0  0  0  0  1  0  1 21  0  0  0  3  0  0]
 [  1  0 109  5  0  2  0  0  0  0  0  0  0  0  0  0  0  0]
 [  0  0  0 209  0  2  0  0  2  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0 99  0  0  0  0  0  0  0  1  3  0  0  9  3]
 [  0  1  5  6  0 186  0  0 19  0  0  0  0  0  0  0  0  0]
 [  0  0  0  0  0  0 209  0  0  1  0  0  0  0  0  0  0 10]
 [  0  0  0  0  0  1 13 64  0  0  0  0  2  0  0  0 26  8]
 [  0  0  0  0  0  0  0  0 110  0  0  0  0  0  4  0  0  0]
 [  1  0  0  0  0  0  0  0  0 108  4  2  0  0  0  0  0  0]
 [  0  0  0  0  0  0  0  0  0  1 213  0  0  0  4  0  0  0]
 [  4  2  0  1  0  0  0  0  7  7  0 198  0  0  0  0  0  0]
 [  0  0  0  0  7  0  2  5  0  0  0  0 63  5  0  1 25  7]
 [  1  0  7  0  0  2  0  2  0  0  0  0  3 184  0  2 10  7]
 [  7  0  1  0  0  0  0  0  5  0  1  0  0  0 198  0  0  0]
 [  0  4  0  0  0  0  0  0  0  0  0  6  0  1  0 118  0  0]
 [  0  0  0  0  1  0  3  8  1  0  0  0  3  0  0  1 96  1]
 [  0  0  0  2  0  0  3  0  0  0  0  2  0  0  0  2  1 208]]
```

Accuracy: 0.8833333333333333

```
In [ ]: from sklearn.metrics import classification_report
print(classification_report(encoded_y_test,y_pred_classes))
```



	precision	recall	f1-score	support
0	0.93	0.89	0.91	219
1	0.92	0.74	0.82	113
2	0.88	0.93	0.90	117
3	0.92	0.98	0.95	213
4	0.93	0.86	0.89	115
5	0.96	0.86	0.91	217
6	0.91	0.95	0.93	220
7	0.81	0.56	0.66	114
8	0.72	0.96	0.82	114
9	0.89	0.94	0.92	115
10	0.97	0.98	0.97	218
11	0.84	0.90	0.87	219
12	0.88	0.55	0.67	115
13	0.95	0.84	0.90	218
14	0.95	0.93	0.94	212
15	0.93	0.91	0.92	129
16	0.57	0.84	0.68	114
17	0.85	0.95	0.90	218
accuracy			0.88	3000
macro avg	0.88	0.87	0.87	3000
weighted avg	0.89	0.88	0.88	3000

```
In [ ]: from sklearn.metrics import precision_recall_fscore_support
res = []
for l in range(18):
    prec, recall, _, _ = precision_recall_fscore_support(np.array(encoded_y_test)==l,
                                                         np.array(y_pred_classes)==l,
                                                         pos_label=True, average=None)
    res.append([l, recall[0], recall[1]])

pd.DataFrame(res, columns = ['class', 'specificity', 'sensitivity'])
```

Out[ ]:

	class	specificity	sensitivity
<b>0</b>	0	0.994966	0.885845
<b>1</b>	1	0.997575	0.743363
<b>2</b>	2	0.994797	0.931624
<b>3</b>	3	0.993541	0.981221
<b>4</b>	4	0.997227	0.860870
<b>5</b>	5	0.997485	0.857143
<b>6</b>	6	0.992446	0.950000
<b>7</b>	7	0.994802	0.561404
<b>8</b>	8	0.985100	0.964912
<b>9</b>	9	0.995494	0.939130
<b>10</b>	10	0.997484	0.977064
<b>11</b>	11	0.986695	0.904110
<b>12</b>	12	0.996880	0.547826
<b>13</b>	13	0.996765	0.844037
<b>14</b>	14	0.996055	0.933962
<b>15</b>	15	0.996865	0.914729
<b>16</b>	16	0.975398	0.842105
<b>17</b>	17	0.987060	0.954128